Pre-requisites

Python

Basics of Spark

Day 1

1) What is Bigdata?
2) Hadoop Yarn Architecture
3) An Introduction to Apache Spark
    a) Spark's Basic Architecture
    b) Spark's Language APIs
    c) Starting Spark
    d) The SparkSession
    e) Data Frames
        i) Partitions
    f) Transformations
        i) Lazy Evaluation
    g) Actions
    h) Spark UI
2) Spark's Toolset
3) Structured API
    a) Data Frames and Datasets
    b) Schemas
    c) Overview of Structured Spark Types
        i) Data Frames Versus Datasets
        ii) Columns
        iii) Rows
        iv) Spark Types

d) Overview of Structured API Execution
    i) Logical Planning
    ii) Physical Planning
    iii) Execution

4) Basic Structured Operations
    a) Schemas
    b) Columns and Expressions
    c) Records and Rows
    d) Data Frame Transformations

5) Executing on Spark
    a) Run job from command line
    b) Go through execution plan / DAG on spark for existing model

6) Using logging

Day 2

8) Details of the Case Study which we will use in our hands-on
9) Aggregations
    a) Aggregation Functions
        i) count
        ii) min and max
        iii) sum
        iv) avg
    b) Grouping
    c) Window Functions
    d) User-Defined Aggregation Functions
10) Joins
    a) Join Expressions
    b) Join Types
    c) Inner Joins
    d) Outer Joins
    e) Left Outer Joins
11) Right Outer Joins
12) Data Sources
    a) The Structure of the Data Sources API
    b) CSV Files
    c) JSON Files
    d) ORC Files
    e) AVRO Files
    f) Parquet Files`
    g) Advanced I/O Concepts
        i) Splitable File Types and Compression
        ii) Reading Data in Parallel

iii) Writing Data in Parallel

iv) Managing File Size

13) Integrating Spark with Hive

14) Spark SQL

a) What Is SQL?

b) How to Run Spark SQL Queries?

c) Catalog

d) Tables

e) Views

f) Select Statements

g) Databases

Day 3

15) Resilient Distributed Datasets (RDDs)
    a) What Are the Low-Level APIs?
    b) About RDDs
        i) Types of RDDs
        ii) When to Use RDDs?
    c) Creating RDDs
    d) Manipulating RDDs
    e) Transformations
        i) distinct
        ii) filter
        iii) map
        iv) sort
        v) Random Splits
    f) Actions
        i) reduce
        ii) count
        iii) first
        iv) max and min
        v) take
    g) Saving Files
        i) saveAsTextFile
        ii) SequenceFiles
        iii) Hadoop Files
16) Caching
17) Persistence
        i) Different types of persistence
        ii) When to use which kind of persistence.

18) Checkpointing
    i) What is checkpointing?
    ii) localCheckpoint
    iii) When and where to use?
19) Distributed Shared Variables
    a) Broadcast Variables
    b) Accumulators
        i) Basic Example
        ii) Custom Accumulators

Day 4

20) Configuring Applications

21) PySpark UDF Introduction

    a) What is UDF?

    b) Why do we need it?

22) Create PySpark UDF (User Defined Function)

    a) Create a Data Frame

    b) Create a Python function

    c) Convert python function to UDF

23) Using UDF with Data Frame

    a) Using UDF with Data Frame select()

    b) Using UDF with Data Frame withColumn()

    c) Registering UDF & Using it on SQL query

24) Introducing Apache Parquet file format

    a) What is Apache Parquet?

    b) Parquet Format vs. CSV

    c) Advantages of Parquet Columnar Storage

    d) Primitive data types in Parquet format

    e) Apache Parquet Spark Example

        i) Spark Write Data Frame to Parquet file format

        ii) Spark Read Parquet file into Data Frame

        iii) Append to existing Parquet file

        iv) Using SQL queries on Parquet

        v) Spark parquet partition – Improving performance

        vi) Spark Read a specific Parquet partition

25) Apache Arrow in PySpark

a) What is Apache Arrow?

b) Apache PyArrow with Apache Spark

c) What is the problem with existing Pandas/Spark conversion without PyArrow?

d) How to use PyArrow in Spark to optimize?

e) Enabling for Conversion to/from Pandas

f) Pandas UDFs

g) Compatibility Setting for PyArrow

h) Converting Pandas Data frame to Apache Arrow Table

i) PyArrow Table to Pandas Data Frame

j) How does the PyArrow enabled conversion work internally?

Day 5

26) Performance Tuning

a) Indirect Performance Enhancements

i) Design Choices

ii) Object Serialization in RDDs

iii) Cluster Configurations

iv) Scheduling

v) Data at Rest

vi) Shuffle Configurations

vii) Memory Pressure and Garbage Collection

b) Direct Performance Enhancements

i) Parallelism

ii) Improved Filtering

iii) Repartitioning and Coalescing

iv) User-Defined Functions (UDFs)

v) Temporary Data Storage (Caching)

vi) Joins

vii) Aggregations

viii) Broadcast Variables

Day 6

29) Machine Learning Basic Concepts
    a) Importing the Libraries
    b) Importing the Dataset
    c) Summary of Object-oriented programming: classes & objects
    d) Missing Data Treatment
    e) Categorical Data
    f) Splitting the Dataset into the Training set and Test set
    g) Feature Scaling
30) Analytics and Machine Learning
    a) What Is Spark's MLlib?
    b) High-Level MLlib Concepts
    c) MLlib in Action
        i) Feature Engineering with Transformers
        ii) Estimators
        iii) Pipelining Our Workflow
        iv) Training and Evaluation
        v) Persisting and Applying Models
31) Preprocessing and Feature Engineering
    a) Formatting Models According to Use Case
    b) Transformers
    c) Estimators for Preprocessing
        i) Transformer Properties
    d) High-Level Transformers
        i) SQL Transformers
        ii) VectorAssembler
    e) Working with Continuous Features
        i) Bucketing

ii) Scaling and Normalization

iii) StandardScaler

f) Working with Categorical Features

i) StringIndexer

ii) Converting Indexed Values Back to Text

iii) Indexing in Vectors

iv) One-Hot Encoding

g) Feature Manipulation

h) Feature Selection

Day 7

32) Pipeline
    a) Building of pipeline
    b) Saving it
    c) Use on a different dataset
33) Classification
    a) Use Cases
    b) Types of Classification
    c) Classification Models in MLlib
    d) Logistic Regression
        i) Logistic Regression Intuition
        ii) Sigmoid Function
        iii) Model Hyperparameters
        iv) Training Parameters
        v) Prediction Parameters
        vi) Example
    e) Decision Trees
        i) Decision Tree Regression Intuition
        ii) Pruning
        iii) Overfitting in Decision Tree
        iv) Entropy
        v) Information Gain
        vi) Model Hyperparameters
        vii) Training Parameters
        viii) Prediction Parameters
34) Regression
    a) Use Cases

b) Regression Models in MLlib

c) Linear Regression

   i) Simple Linear Regression Intuition

   ii) RMSE

   iii) Model Hyperparameters

   iv) Training Parameters

   v) Example

   vi) Training Summary

d) Decision Trees

   i) Model Hyperparameters

   ii) Training Parameters

   iii) Example